



Computer Vision

Lecture 2: Edge Detection

30.10.2024

Manuel Heurich, Tim Landgraf




List of Topics

<ul style="list-style-type: none">• Introduction• Convolution, Edge Detection• Color Histograms, HoG	Computer Vision Fundamentals
<ul style="list-style-type: none">• Optic Flow• Hough Transform• SIFT / SURF	Conventional Computer Vision
<ul style="list-style-type: none">• Introduction to Neural Information Processing• Convolutional Neural Networks• Image Classification, Object Detection• Vision Transformers, ConvMixer• Semantic Segmentation• Pose Estimation• Recurrent Neural Networks, Image Captioning	Deep Learning (Supervised)
<ul style="list-style-type: none">• Generative Models• Unsupervised feature extraction	Unsupervised Learning



Blob to Center of Mass



Collect all foreground pixels' positions

$$B = \{(x_1, y_1)^T, (x_2, y_2)^T, \dots, (x_n, y_n)^T\}$$

Compute average

$$c = \frac{\sum_i \begin{pmatrix} x_i \\ y_i \end{pmatrix}}{n}$$

$$c = \frac{\begin{pmatrix} 18 \\ 20 \end{pmatrix}}{7} = \begin{pmatrix} 2,5 \\ 2,9 \end{pmatrix}$$





Tracking

- Given: set of detections

$$D^{(t)} = \{p_i\}; i = 1 \dots N, ; p_i = \begin{pmatrix} x_i^{(t)} \\ y_i^{(t)} \\ \dots \end{pmatrix}$$


at time t and $t + 1$

- Find: the mapping $M: D_i^{(t)} \rightarrow D_j^{(t+1)}$ that retains the identity of the objects
- Note: in general, objects may (re, dis-)appear any time






Simplest tracking algorithm






Simplest tracking algorithm





Simplest tracking algorithm




Algorithm:

```
C = get_candidates()
for last detection d in D
    c = get_closest_candidate(d, C)
    M.add(d, c)
    C.remove(c)
```



Simplest tracking algorithm




When does this **fail**?

Is the results **invariant to the order** of our last detections?

Will those problems become relevant when tracking a **single object**?



Simplest tracking algorithm



What can we **change in our recording setup** to improve tracking quality?




Ex. 1.3: track colored objects

- use your color detection and connected components algorithm
- implement simplest tracking algorithm
- draw history of all previous points on frame





Today: EDGES




White eagle on white background



Edges are robust features


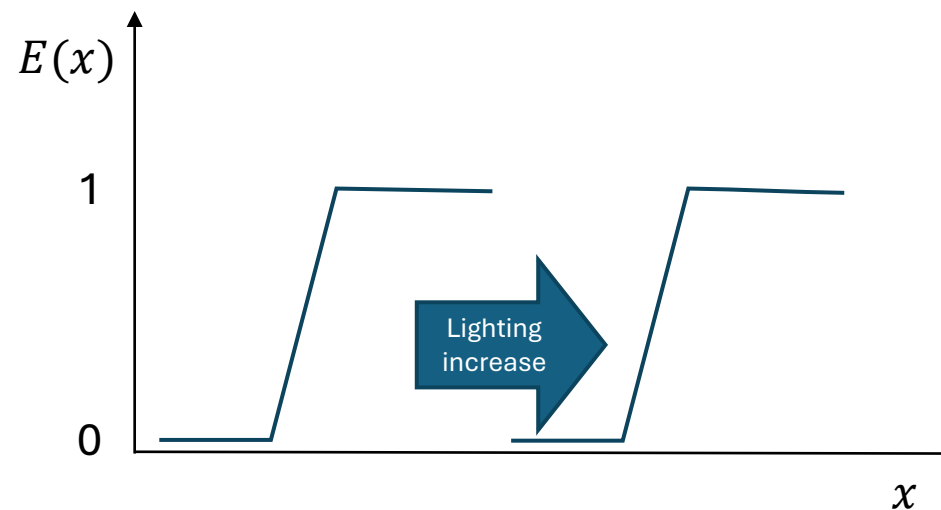
- global lighting changes introduce brightness shifts, but local brightness differences remain stable





Edges are robust features



- global lighting changes introduce brightness shifts, but local brightness differences remain stable
- edge image = binary image
- detect local brightness changes!









Properties of brightness changes




High slope




Low slope



Properties of brightness changes



No slope in x direction



High slope in y direction



Image Derivative

$$I'(x) = \frac{I(x + \Delta) - I(x)}{\Delta}$$














Image Derivative

$$I'(x) = \frac{I(x + \Delta) - I(x)}{\Delta}$$






Computer Vision 101: Convolution

$$(I * K)_{x,y} = \sum_{-W}^W \sum_{-H}^H I_{x-w,y-h} K_{w,h}$$
$$= I_{N(x,y)}^T K$$

Notion 1: weighted sum






Computer Vision 101: Convolution

$$(I * K)_{x,y} = \sum_{-W}^W \sum_{-H}^H I_{x-w,y-h} K_{w,h}$$
$$= I_{N(x,y)}^T K$$

Notion 2: pattern matching via dot product






Convolution

$$K = \begin{bmatrix} [1, 1, 1, 1, 1], \\ [1, 1, 1, 1, 1], \\ [1, 1, 1, 1, 1] \end{bmatrix}$$






„Detect“ edges with Convolution





„Detect“ edges with Convolution


$$K = \begin{bmatrix} [-1, & 0, & 1], \\ [-1, & 0, & 1], \\ [-1, & 0, & 1] \end{bmatrix}$$





„Detect“ edges with Convolution

$$K = \begin{bmatrix} [-1, 0, 1], \\ [-1, 0, 1], \\ [-1, 0, 1] \end{bmatrix}^T$$





Sobel Kernel, Sobel Filter

$$S_x = \begin{bmatrix} [-1, & 0, & 1], \\ [-2, & 0, & 2], \\ [-1, & 0, & 1] \end{bmatrix}$$

Convolve for derivative in x-direction

$$G_x = (I * S_x)$$

$$S_y = \begin{bmatrix} [-1, & -2, & -1], \\ [0, & 0, & 0], \\ [1, & 2, & 1] \end{bmatrix}$$

Convolve for derivative in y-direction

$$G_y = (I * S_y)$$




The Gradient

- The vector of partial image derivatives

$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y)$$


is called the gradient.

- The gradient is pointing to the direction of maximum slope and is orthogonal to the edge's direction






Computing the Image Gradient



$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} ? \\ ? \end{pmatrix}$$




Computing the Image Gradient




$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$



Computing the Image Gradient




$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$



$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} ? \\ ? \end{pmatrix}$$




Computing the Image Gradient





$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$

Use Sobel filter to estimate the image gradient



$$\nabla I(x, y) = \begin{pmatrix} \frac{\delta I}{\delta x} \\ \frac{\delta I}{\delta y} \end{pmatrix} (x, y) = \begin{pmatrix} ? \\ ? \end{pmatrix}$$

$$\nabla I(x, y) = \begin{pmatrix} G_x \\ G_y \end{pmatrix} (x, y)$$







Canny's Edge Detection

[Cited by 29138](#)

IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. PAMI-8, NO. 6, NOVEMBER 1986

A Computational Approach to Edge Detection

JOHN CANNY, MEMBER, IEEE

Abstract—This paper describes a computational approach to edge detection. The success of the approach depends on the definition of a *comprehensive* set of goals for the computation of edge points. These goals must be precise enough to delimit the desired behavior of the detector while making minimal assumptions about the form of the solution. We define detection and localization criteria for a class of edges, and present mathematical forms for these criteria as functionals on the operator impulse response. A third criterion is then added to ensure that the detector has only one response to a single edge. We use the criteria in numerical optimization to derive detectors for several common image features, including step edges. On specializing the analysis to step edges, we find that there is a natural uncertainty principle between detection and localization performance, which are the two main goals. With this principle we derive a single operator shape which is optimal at any scale. The optimal detector has a simple approximate implementation in which edges are marked at maxima in gradient magnitude of a Gaussian-smoothed image. We extend this simple detector using operators of several widths to cope with different signal-to-noise ratios in the images. We present a general method, called feature con-

detector as input to a program which could isolate simple geometric solids. More recently the model-based vision system ACRONYM [3] used an edge detector as the front end to a sophisticated recognition program. Shape from motion [29], [13] can be used to infer the structure of three-dimensional objects from the motion of edge contours or edge points in the image plane. Several modern theories of stereopsis assume that images are preprocessed by an edge detector before matching is done [19], [20]. Beattie [1] describes an edge-based labeling scheme for low-level image understanding. Finally, some novel methods have been suggested for the extraction of three-dimensional information from image contours, namely shape from contour [27] and shape from texture [31].

In all of these examples there are common criteria relevant to edge detector performance. The first and most


67





Requirements for a good edge detector

- detection of edge with **low error rate** (low false positives and false negative rates)
- edge point should be **localized** accurately on the center of the edge
- **no double detections!**





Outline of Canny's Edge Detection

- **convolve** with Gaussian to smooth the image and remove noise
- **convolve** again with Sobel to compute the image gradient
- apply **non-maximum suppression**
- use **two thresholds**: 1) for strong and 2) weak edges
- **hysteresis**: track edge along edge direction and suppress weak edges not connected to strong edges



Gaussian Smoothing

$$K = \begin{bmatrix} ? & ? & ? \end{bmatrix}$$

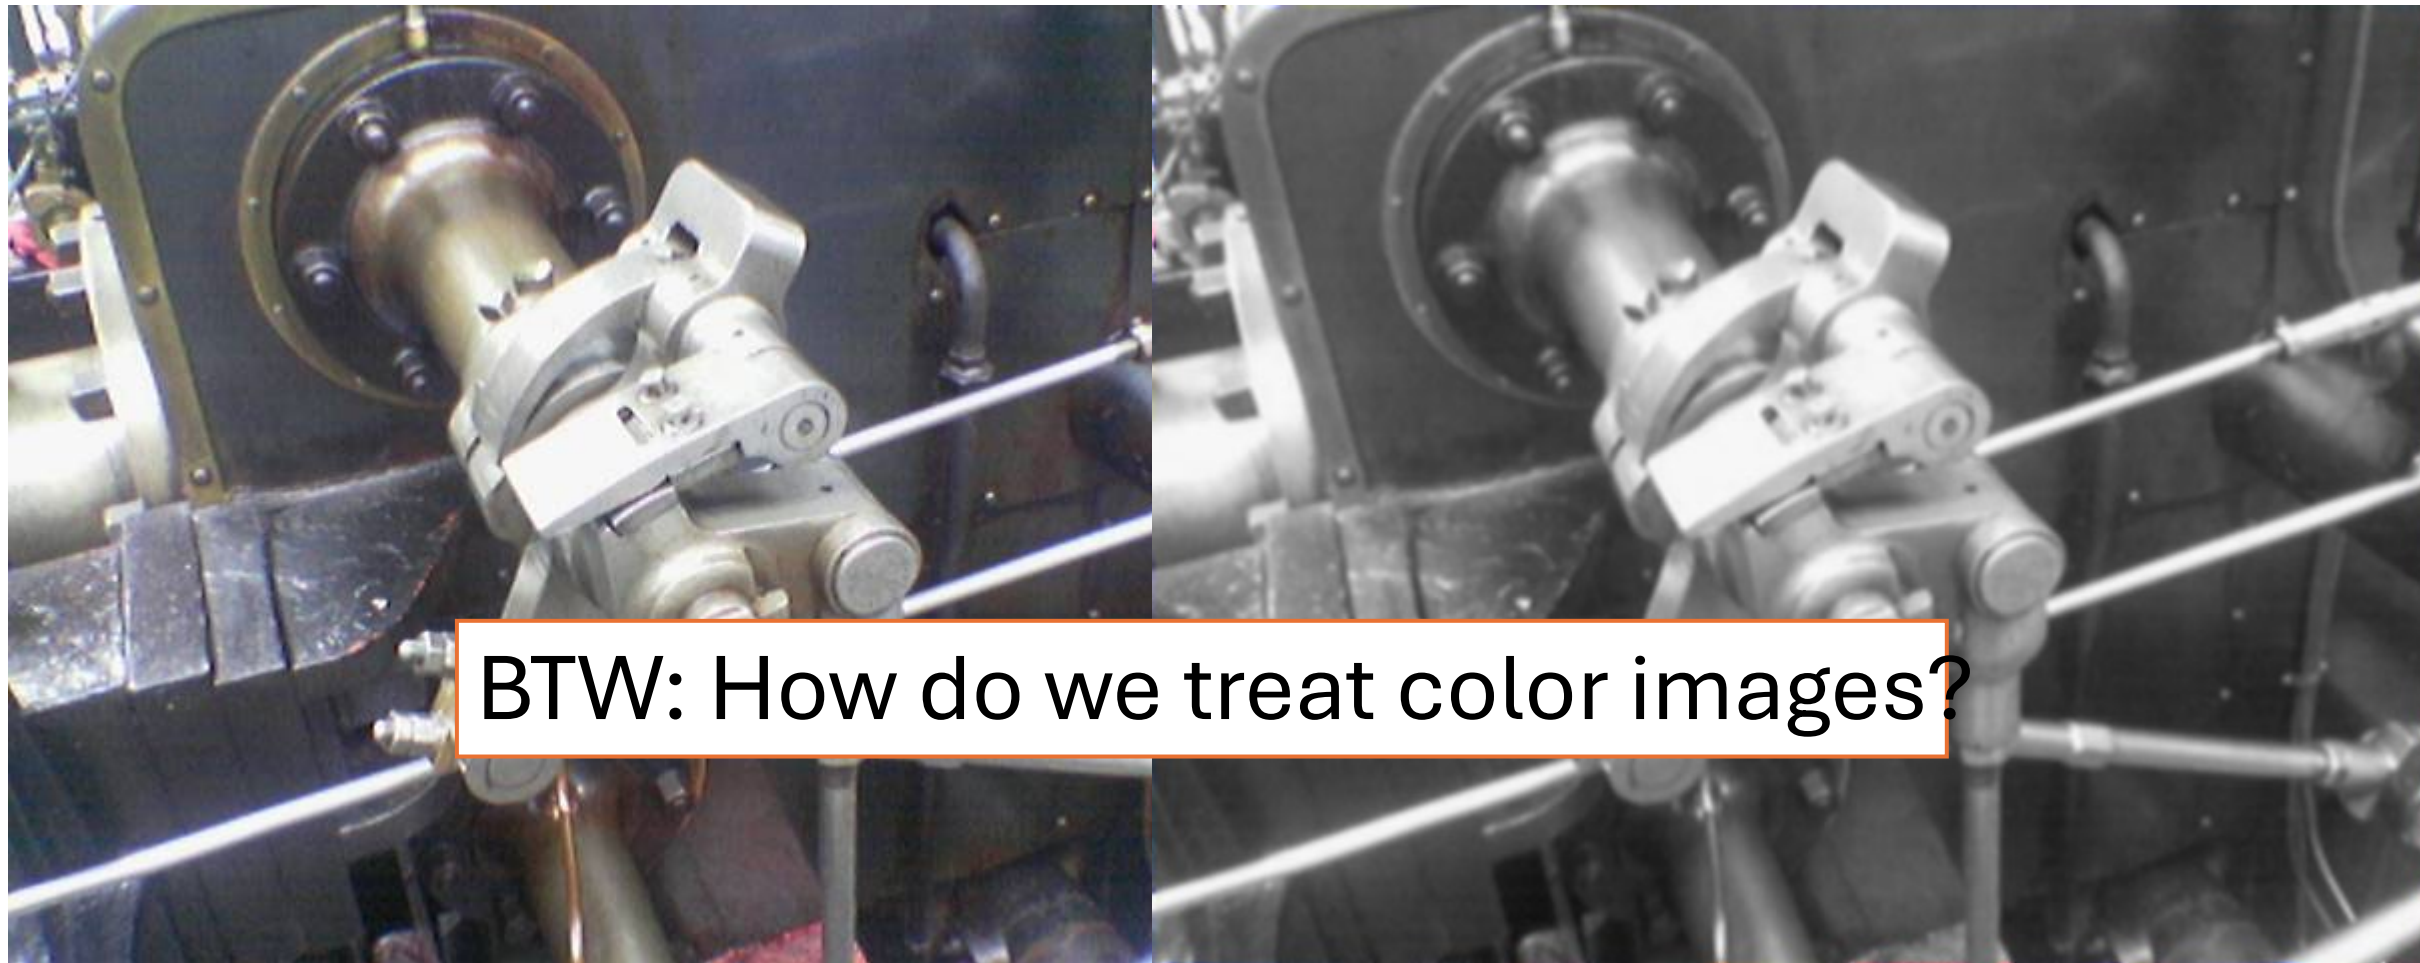




Image gradient magnitude and orientation

- Convolve with Sobel: $G_x = I * S_x$ and $G_y = I * S_y$ ✓
- The **gradient magnitude** is the length of the gradient vector
- atan2 returns the edge orientation between 0 and 2π
- θ is then rounded to eight discrete directions (corresponding to the 8-neighborhood)



$$G = \sqrt{G_x^2 + G_y^2}$$

$$\Theta = \text{atan2}(G_y, G_x)$$



Non-maximum suppression: edge thinning

- Scan over all pixels, compare edge magnitude with that of its neighbors (as indicated by gradient orientation)
- Suppress if not greater





Non-maximum suppression: edge thinning


- Scan over all pixels, compare edge magnitude with that of its neighbors (as indicated by gradient orientation)
- Suppress if not greater
- **Q:** For which gradient direction would the center pixel not be suppressed?

5	10	10
11	10	6
10	10	3



Weak or strong edges?

- Two thresholds
- one defining strong edges, one for weak edges, else: no edge!
- Strong pixels: should be in the result set, weak pixels only if they are connected to strong edges!





Effect of Low/High Thresholds

(using relative thresholds here:

0.4 means 40% of the maximum gradient magnitude)

$\sigma = 1, l = 0.4, h = 0.4$


$\sigma = 1, l = 0.4, h = 0.6$



$\sigma = 1, l = 0.2, h = 0.6$



$\sigma = 1, l = 0.2, h = 0.4$




$\sigma = 1, l = 0.2, h = 0.2$





Effect of smoothing



$\sigma = 1$

$\sigma = 3$



Summary (aka “what will likely show up in the finals”)

- “Edges” = binary concept
 - Which properties can we use to detect them?
 - Requirements for “good” edge detectors
- Image derivatives via convolution with various kernels (e.g. Sobel)
 - What is convolution? What is it used for?
 - What is the Sobel kernel used for?
- Edge detection with Canny’s edge detector
 - Describe the steps
 - What are expected effects when changing parameters such as σ , θ_h , θ_l , etc?