



Computer Vision

Lecture 4: Histograms

13.11.2024

Manuel Heurich, Tim Landgraf



List of Topics

1.	Introduction (Color)	Computer Vision Fundamentals
2.	Edge Detectors	
3.	Hough Transform	Conventional Computer Vision
4.	Histogram-based Detection and Tracking	
5.	Optic flow	
6.	SIFT / SURF	
7.	Introduction to Neural Information Processing	Deep Learning (Supervised)
8.	Convolutional Neural Networks	
9.	Image Classification, Object Detection	
10.	Semantic Segmentation	
11.	Pose Estimation	
12.	Style Transfer	
13.	Recurrent Neural Networks, Image Captioning	
14.	Generative Models	Unsupervised Learning
15.	Unsupervised feature extraction	



(Bradski, 1998)

Cited by 2052

Computer Vision Face Tracking For Use in a Perceptual User Interface

Gary R. Bradski, Microcomputer Research Lab, Santa Clara, CA, Intel Corporation



Index words: computer vision, face tracking, mean shift algorithm, perceptual user interface, 3D graphics interface

Abstract

As a first step towards a perceptual user interface, a computer vision color tracking algorithm is developed and applied towards tracking human faces. Computer vision algorithms that are intended to form part of a perceptual user interface must be fast and efficient. They must be able to track in real time yet not absorb a major share of computational resources; other tasks must be able to run


this effort: a 4-degree of freedom color object tracker and its application to flesh-tone-based face tracking.

Computer vision face tracking is an active and developing field, yet the face trackers that have been developed are not sufficient for our needs. Elaborate methods such as tracking contours with snakes [[10][12][13]], using Eigenspace matching techniques [14], maintaining large sets of statistical hypotheses [15], or convolving images






Today: Histograms as Object Features



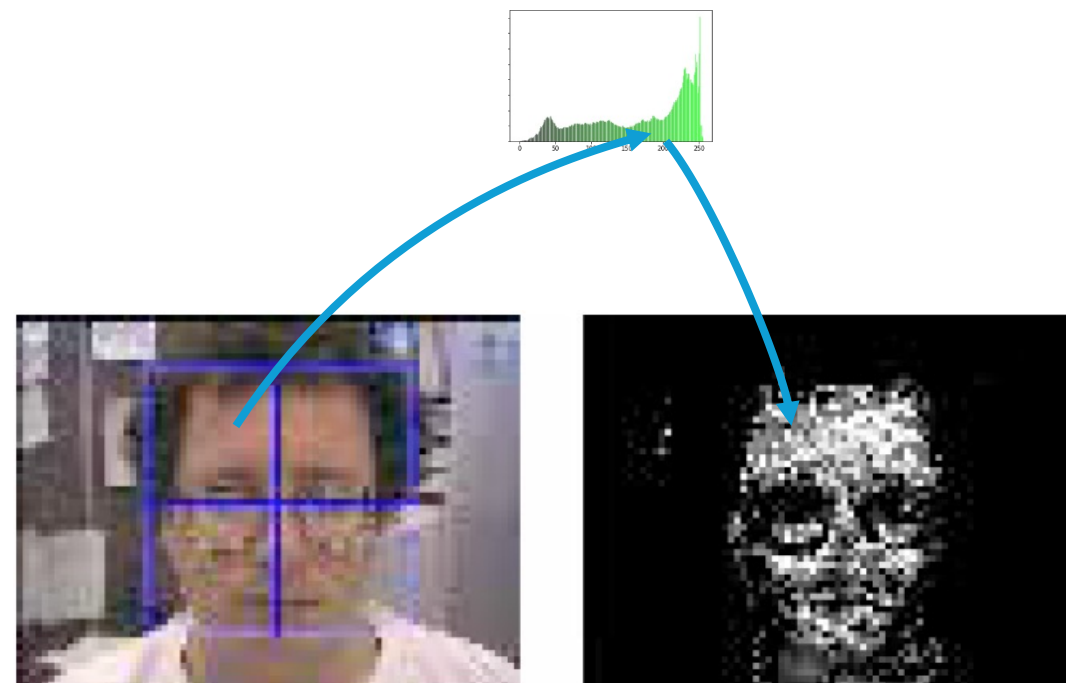
Lecture #1: color sub-space defines all colors that belong to an object



Today: Histograms as Object Features



Lecture #1: color sub-space defines all colors that belong to an object



Today: color histogram defines likelihood of colors to belong to an object





Color Histograms

Image



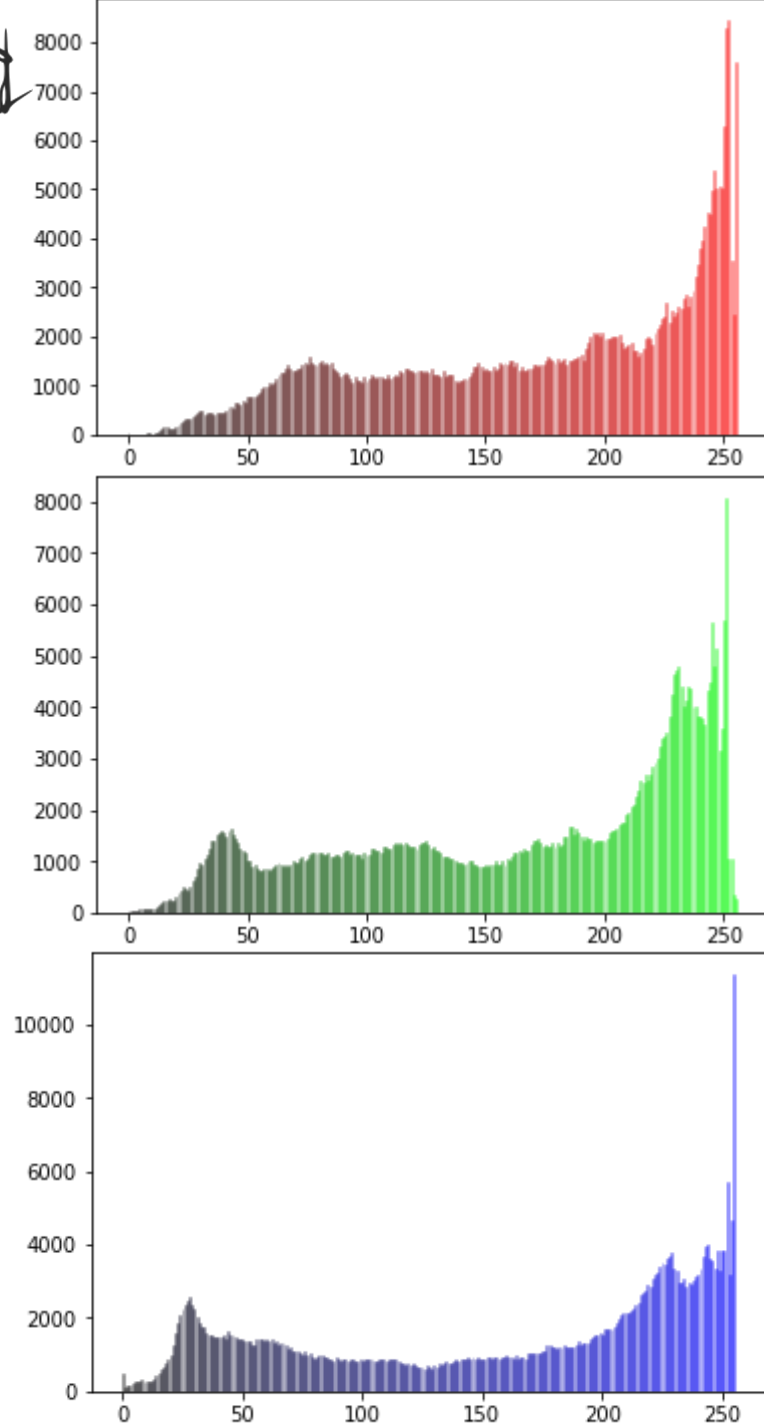
Histogram







3D-Histogram \neq 3 Histograms!






3D-Histogram, actual

Index into histogram with rgb-triplet


Every cell contains count of that specific triplet

Q: What is the problem with a 3D histogram?





HSV, a better color space?




HSV representation models how colors appear under light

- **Hue**: angular dimension, represents color tone from red (0°), green (120°) to blue (240°)
- **Saturation**: S_{\max} is purest color, mix with white \rightarrow decrease S
- **Value**: shine light on color, V_{\max} yields most intense color



Hue histogram







Histogram as probability map


The distribution of color hues

- reflects how likely a specific hue is within the object boundaries
- is an **object descriptor**
- given a pixel p with hue h , how likely belongs p to the object?






Track the position of the distribution: Mean Shift



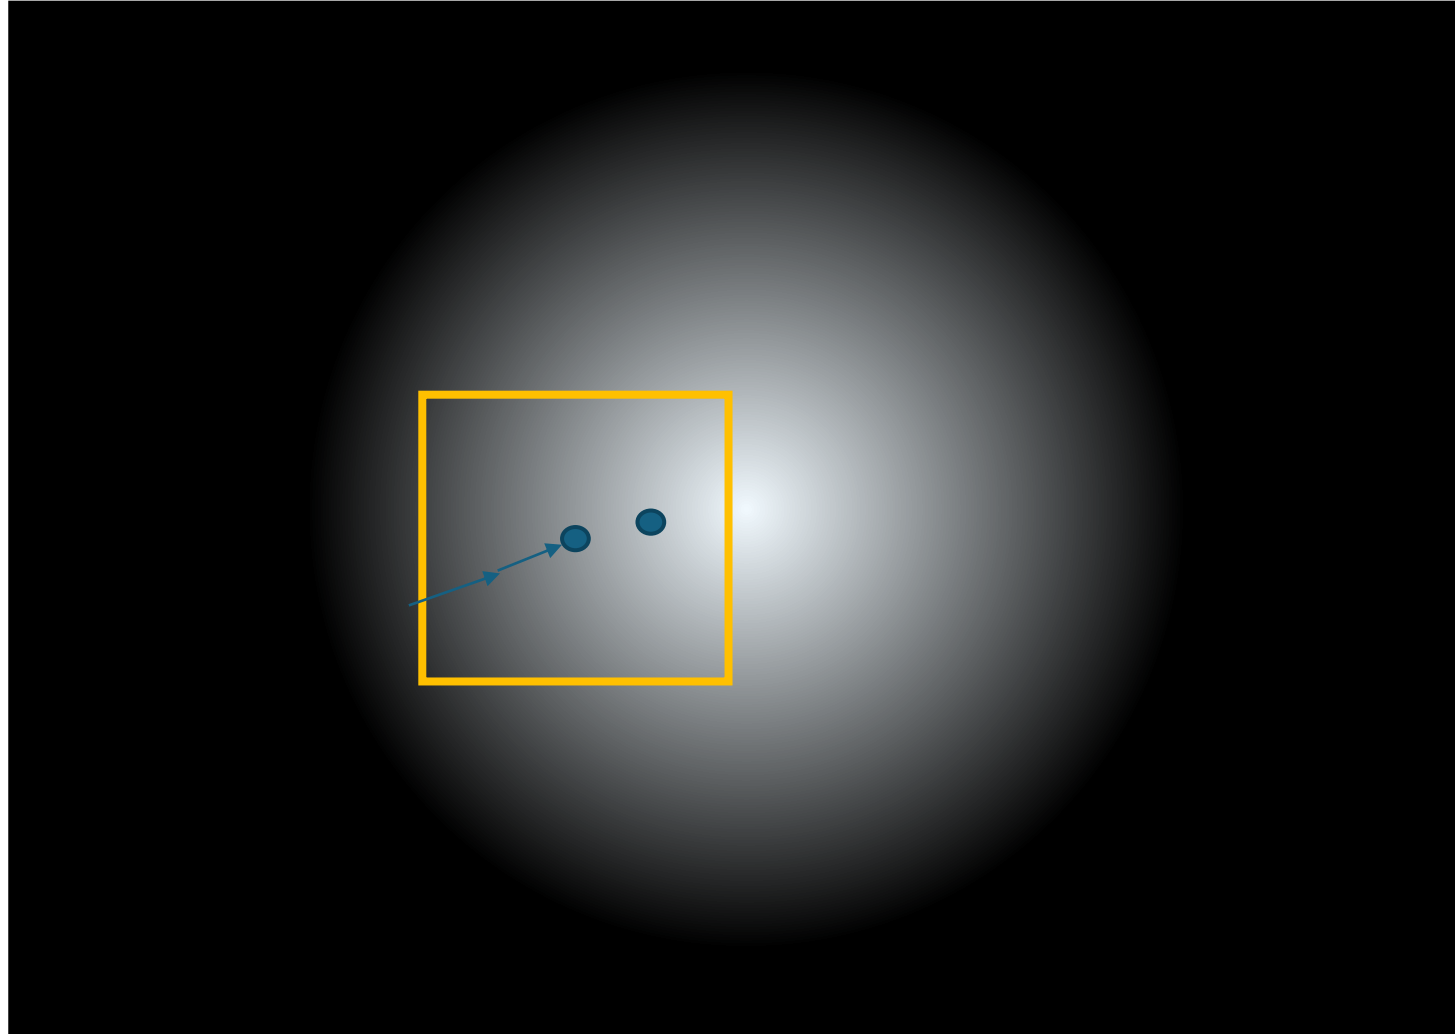


Track the position of the distribution: Mean Shift






Track the position of the distribution: Mean Shift





Track the position of the distribution: Mean Shift






Mean Shift: using image moments

$$M_{ij} = \sum_x \sum_y x^i y^j I(x, y)$$

$$M_{00} = \sum_x \sum_y I(x, y) \quad \text{zeroth moment}$$

$$M_{10} = \sum_x \sum_y x I(x, y) \quad M_{01} = \sum_x \sum_y y I(x, y) \quad \text{first moments}$$



$$x_c = \frac{M_{10}}{M_{00}} \quad y_c = \frac{M_{01}}{M_{00}}$$

center of mass



Ex. 3.1: track colored objects with mean shift

- read the CAMSHIFT paper
- use an initial search window to extract an object's hue histogram
- implement the probability image given the histogram descriptor
- track the mode of the color distribution with mean shift



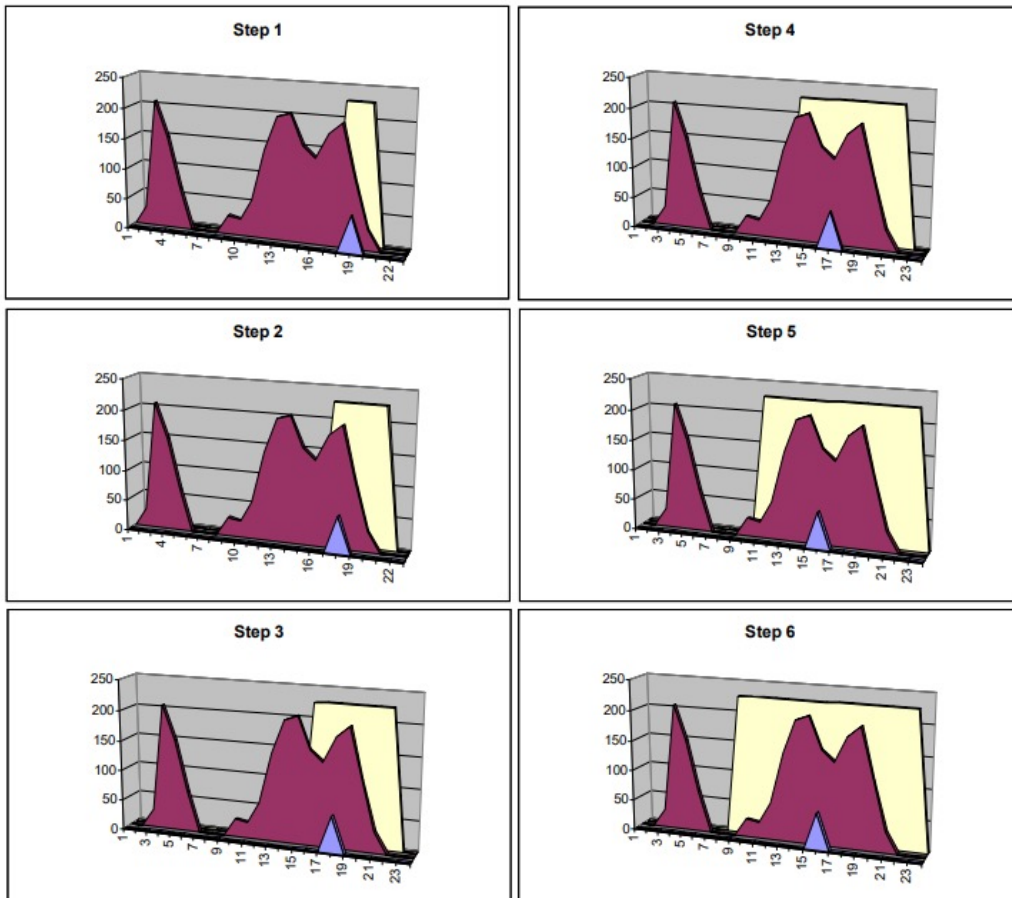


What if we lose track of the object?





CAMSHIFT: adaptive window size




Update search window size
depending on probability
mass in window:

$$s = 2 \sqrt{\frac{M_{00}}{256}}$$

Q: What happens with non-zero noise levels?



CAMSHIFT: Track the orientation of the object




<https://www.youtube.com/watch?v=iBOLbs8i7Og>



CAMSHIFT: object orientation

$$\theta = \frac{\operatorname{atan}\left(\frac{2\left(\frac{M_{11}}{M_{00}} - x_c y_c\right)}{\left(\frac{M_{20}}{M_{00}} - x_c^2\right) - \left(\frac{M_{02}}{M_{00}} - y_c^2\right)}\right)}{2}$$

$$M_{ij} = \sum_x \sum_y x^i y^j I(x, y)$$



θ is the angle of the first principal component of the probability distribution!

Read more here:

https://en.wikipedia.org/wiki/Image_moment

https://en.wikipedia.org/wiki/Principal_component_analysis

Q: Why can't we just use the eigenvector of the covariance matrix?



Ex. 3.2: track colored objects with CAMSHIFT

- implement the automatic determination of the search window size
- implement the extraction of object orientation and ellipse axes
- draw a respective ellipse on the image(s)





Summary: Color-based object detection

- color *can be* a stable feature, when using the right color space (HSV rather than RGB)
- color does not encode *structural* image features
- color histograms can serve as **object descriptor** and **predictor** of a pixels object membership (see probability image in CAMSHIFT)
- even in HSV color may vary, depending on various factors





HOG: Histograms of Oriented Gradients

(Dalal & Triggs, 2005)

[Cited by 22193](#)

Histograms of Oriented Gradients for Human Detection

Navneet Dalal and Bill Triggs

INRIA Rhône-Alpes, 655 avenue de l'Europe, Montbonnot 38334, France

{Navneet.Dalal,Bill.Triggs}@inrialpes.fr, <http://lear.inrialpes.fr>

Abstract

We study the question of feature sets for robust visual object recognition, adopting linear SVM based human detection as a test case. After reviewing existing edge and gradient based descriptors, we show experimentally that grids of Histograms of Oriented Gradient (HOG) descriptors significantly outperform existing feature sets for human detection. We study the influence of each stage of the computation on performance, concluding that fine-scale gradients, fine orientation binning, relatively coarse spatial binning, and high-quality local contrast normalization in overlapping de-



We briefly discuss previous work on human detection in §2, give an overview of our method §3, describe our data sets in §4 and give a detailed description and experimental evaluation of each stage of the process in §5–6. The main conclusions are summarized in §7.


2 Previous Work

There is an extensive literature on object detection, but here we mention just a few relevant papers on human detection [18, 17, 22, 16, 20]. See [6] for a survey. Papageorgiou *et al* [18] describe a pedestrian detector based on a polynomial SVM using modified Haar-like features on input detection with




Remember lecture 2: Convolution, Gradients and Edge Detection






HOG example:
pedestrian
detection (and
localization)

Detection (of pedestrians) =
classification of ROIs





Feature space

Detection (of pedestrians) =
classification of ROIs



Feature space

Detection (of pedestrians) =
classification of ROIs




Feature space



ROI from Blocks, Blocks from Cells

Cells (8 x 8 pixels)


Blocks (2 x 2 cells, stride : 1 cell)





Creating Gradient Histograms


- 9 bins for 180° (Q: why not 360° ?)





Creating Gradient Histograms


- 9 bins for 180° (Q: why not 360° ?)
- For each pixel in cell C_i compute ($8 \times 8 = 64$) gradient vectors
- *Don't count, add* gradient magnitudes - “voting by bilinear interpolation”





Block descriptor

- For each block of 2 x 2 cells, concatenate the 4 histograms to a block descriptor **b**
- Normalize **b**:
$$b \leftarrow \frac{b}{\sqrt{\|b\|^2 + \epsilon}}$$
- Since each cell is covered by four blocks, each histogram is represented four times with four different normalizations (except for cells at the border of the ROI)





HOG descriptor

- All block descriptors in a ROI are concatenated into a HOG feature vector h

- Normalize h :
$$h \leftarrow \frac{h}{\sqrt{\|h\|^2 + \epsilon}}$$


- Clip every entry :
$$h_n \leftarrow \min(h_n, \tau)$$

- Normalize again :
$$h \leftarrow \frac{h}{\sqrt{\|h\|^2 + \epsilon}}$$

- With ROI 128 x 64, we have 16 x 8 cells and 15 x 7 blocks, resulting in a 3780-dimensional feature vector



h now needs to be classified



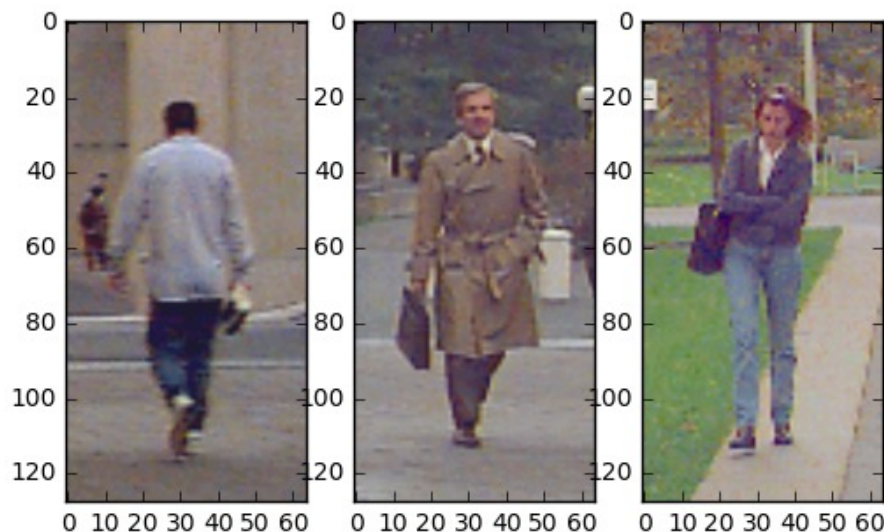
- Binary classification (pedestrian vs. Rest)
- The original work used a SVM
- Others have used random forests, gaussian mixtures, etc.
- The classifier has to be trained with labeled data





Ex. 4.3: track objects with HOG

- read the paper
- implement and test gradient histograms
- implement HOG descriptor
- detect pedestrians (using kNN, SVM, Logistic Regression)





Summary

- HOG features are based on histograms of gradient orientations (of so-called cells), concatenated over blocks to create a block descriptor, further concatenated to form the feature vector
- Special type of histogram: we sum up the gradient magnitudes rather than count the gradients that fall in a histogram bin
- HOG can be used to detect not just pedestrians (given a respective dataset)
- HOG is the feature, the classifier can be anything you like